

Assignment 3- Parallel Molecular Dynamics

CSCI 596: Scientific Computing & Visualization

Anup V Kanale

October 4, 2017

This purpose of this assignment is run molecular dynamics on multiple processors to understand asynchronous message passing and "in-situ" data analysis.

1 Asynchronous Messages

In this section, we use `MPI_Irecv()` and `MPI_Send()` with an `MPI_Wait()` in the program `pmd.c`. The asynchronous messages make the deadlock-avoidance scheme unnecessary, and thus there is no need to use different orders of send and receive calls for even and odd-parity processes.

The new asynchronous implementation takes $5.13e-01$ seconds, whereas the original implementation takes $6.28e-01$ seconds. This observed speedup is $\approx 20\%$. The major modifications were made in the functions `atom_copy` and `atom_move`. To keep this report short, only these parts are shown in Appendix A.1.

2 Communicators

In this section, we basically allocate half the processors for the computations, and the other half for analysing the data from the computations. Specifically, following the lecture note on "In situ analysis of molecular dynamics simulation data using communicators", `pmd.c` is modified such that as many number of processes as that for MD simulations are spawned to calculate the probability density function (PDF) for the atomic velocity.

The plot of the probability density function is as shown below: The major modification was in




Figure 1: Probability density function

the `main` function, `calc_pv()` was provided. Only the modified portion is shown in appendix A.2 to keep this report short.

A Appendix

A.1 Task I

```

1  /*-----*/
2  void atom_copy() {
3  /*-----
4  Exchanges boundary-atom coordinates among neighbor nodes: Makes
5  boundary-atom list, LSB, then sends & receives boundary atoms.
6  -----*/
7  int kd,kdd,i,ku,inode,nsd,nrc,a;
8  int nbnew = 0; /* # of "received" boundary atoms */
9  double com1;
10
11 /* Main loop over x, y & z directions
12    starts-----*/
13  for (kd=0; kd<3; kd++) {
14
15      /* Make a boundary-atom list,
16         LSB-----*/
17
18      /* Reset the # of to-be-copied atoms for lower&higher directions */
19      for (kdd=0; kdd<2; kdd++) lsb[2*kd+kdd][0] = 0;
20
21      /* Scan all the residents & copies to identify boundary atoms */
22      for (i=0; i<n+nbnew; i++) {
23          for (kdd=0; kdd<2; kdd++) {
24              ku = 2*kd+kdd; /* Neighbor ID */
25              /* Add an atom to the boundary-atom list, LSB, for neighbor ku
26                 according to bit-condition function, bbd */
27              if (bbd(r[i],ku)) lsb[ku][++(lsb[ku][0])] = i;
28          }
29      }
30
31      /* Message
32         passing-----*/
33
34      com1=MPI_Wtime(); /* To calculate the communication time */
35
36      /* Loop over the lower & higher directions */
37      for (kdd=0; kdd<2; kdd++) {
38          inode = nn[ku=2*kd+kdd]; /* Neighbor node ID */

```

```

38
39      /* Send & receive the # of boundary
        atoms—————*/
40
41      nsd = lsb[ku][0]; /* # of atoms to be sent */
42
43      //—START
        MODIFIED—————
44
45      MPI_Irecv(&nrc,1,MPI_INT,MPI_ANY_SOURCE,10,MPI_COMM_WORLD,&Request);
46      MPI_Send(&nsd,1,MPI_INT,inode,10,MPI_COMM_WORLD);
47      MPI_Wait(&Request, &status);
48
49      /* Send & receive information on boundary
        atoms—————*/
50      MPI_Irecv(dbufr,3*nrc,MPI_DOUBLE,MPI_ANY_SOURCE,20,MPI_COMM_WORLD,&Request);
51
52      /* Message buffering */
53      for (i=1; i<=nsd; i++)
54          for (a=0; a<3; a++) /* Shift the coordinate origin */
55              dbuf[3*(i-1)+a] = r[lsb[ku][i]][a]-sv[ku][a];
56
57      MPI_Send(dbuf,3*nsd,MPI_DOUBLE,inode,20,MPI_COMM_WORLD);
58      MPI_Wait(&Request, &status);
59      //—END
        MODIFIED—————
60
61
62      /* Message storing */
63      for (i=0; i<nrc; i++)
64          for (a=0; a<3; a++) r[n+nbnew+i][a] = dbufr[3*i+a];
65
66      /* Increment the # of received boundary atoms */
67      nbnew = nbnew+nrc;
68
69      /* Internode synchronization */
70      MPI_Barrier(MPI_COMM_WORLD);
71
72      } /* Endfor lower & higher directions, kdd */
73
74      comt += MPI_Wtime()-com1; /* Update communication time, COMT */
75
76      } /* Endfor x, y & z directions, kd */
77
78      /* Main loop over x, y & z directions
        ends—————*/
79
80      /* Update the # of received boundary atoms */

```

```

81  nb = nbnew;
82 }
83
84
85 /*-----*/
86 void atom_move() {
87 /*-----
88 Sends moved-out atoms to neighbor nodes and receives moved-in atoms
89 from neighbor nodes. Called with n, r[0:n-1] & rv[0:n-1], atom_move
90 returns a new n' together with r[0:n'-1] & rv[0:n'-1].
91 -----*/
92
93 /* Local
   variables-----
94
95 mvque[6][NBMAX]: mvque[ku][0] is the # of to-be-moved atoms to neighbor
96 ku; MVQUE[ku][k>0] is the atom ID, used in r, of the k-th atom to be
97 moved.
98 -----*/
99 int mvque[6][NBMAX];
100 int newim = 0; /* # of new immigrants */
101 int ku,kd,i,kdd,kul,kuh,inode,ipt,a,nsd,nrc;
102 double com1;
103
104 /* Reset the # of to-be-moved atoms, MVQUE[][0] */
105 for (ku=0; ku<6; ku++) mvque[ku][0] = 0;
106
107 /* Main loop over x, y & z directions
   starts-----*/
108
109 for (kd=0; kd<3; kd++) {
110
111     /* Make a moved-atom list,
       mvque-----*/
112
113     /* Scan all the residents & immigrants to list moved-out atoms */
114     for (i=0; i<n+newim; i++) {
115         kul = 2*kd ; /* Neighbor ID */
116         kuh = 2*kd+1;
117         /* Register a to-be-copied atom in mvque[kul/kuh][] */
118         if (r[i][0] > MOVED_OUT) { /* Don't scan moved-out atoms */
119             /* Move to the lower direction */
120             if (bmv(r[i],kul)) mvque[kul][++(mvque[kul][0])] = i;
121             /* Move to the higher direction */
122             else if (bmv(r[i],kuh)) mvque[kuh][++(mvque[kuh][0])] = i;
123         }
124     }
125

```

```

126      /* Message passing with neighbor
        nodes-----*/
127
128      com1 = MPI_Wtime();
129
130      /* Loop over the lower & higher
        directions-----*/
131
132      for (kdd=0; kdd<2; kdd++) {
133
134          inode = nn[ku=2*kd+kdd]; /* Neighbor node ID */
135
136          /* Send atom-number
            information-----*/
137
138          nsd = mvque[ku][0]; /* # of atoms to-be-sent */
139
140      //---START
        MODIFIED-----
141          /* Send & receive information on boundary
            atoms-----*/
142          MPI_Irecv(&nrc,1,MPI_INT,MPI_ANY_SOURCE,110,MPI_COMM_WORLD,&Request);
143          MPI_Send(&nsd,1,MPI_INT,inode,110,MPI_COMM_WORLD);
144          MPI_Wait(&Request, &status);
145
146          MPI_Irecv(dbufr,6*nrc,MPI_DOUBLE,MPI_ANY_SOURCE,120,MPI_COMM_WORLD,&Request);
147
148          /* Message buffering */
149          for (i=1; i<=nsd; i++)
150              for (a=0; a<3; a++) {
151                  /* Shift the coordinate origin */
152                  dbuf[6*(i-1)+a] = r[mvque[ku][i]][a]-sv[ku][a];
153                  dbuf[6*(i-1)+3+a] = rv[mvque[ku][i]][a];
154                  r[mvque[ku][i]][0] = MOVED_OUT; /* Mark the moved-out atom */
155              }
156
157          MPI_Send(dbuf,6*nsd,MPI_DOUBLE,inode,120,MPI_COMM_WORLD);
158          MPI_Wait(&Request, &status);
159      //---END
        MODIFIED-----
160
161
162      /* Message storing */
163      for (i=0; i<nrc; i++)
164          for (a=0; a<3; a++) {
165              r[n+newim+i][a] = dbufr[6*i+a];
166              rv[n+newim+i][a] = dbufr[6*i+3+a];
167          }

```

```

168
169     /* Increment the # of new immigrants */
170     newim = newim+nrc;
171
172     /* Internode synchronization */
173     MPI_Barrier(MPI_COMM_WORLD);
174
175     } /* Endfor lower & higher directions , kdd */
176
177     comt=comt+MPI_Wtime()-com1;
178
179     } /* Endfor x, y & z directions , kd */
180
181     /* Main loop over x, y & z directions
182        ends—————*/
183
184     /* Compress resident arrays including new immigrants */
185     ipt = 0;
186     for (i=0; i<n+newim; i++) {
187         if (r[i][0] > MOVED_OUT) {
188             for (a=0; a<3; a++) {
189                 r[ipt][a] = r[i][a];
190                 rv[ipt][a] = rv[i][a];
191             }
192             ++ipt;
193         }
194     }
195
196     /* Update the compressed # of resident atoms */
197     n = ipt;
198 }
199
200 /*—————*/
201 Bit condition functions:
202
203 1. bdd(ri,ku) is .true. if coordinate ri[3] is in the boundary to
204   neighbor ku.
205 2. bmv(ri,ku) is .true. if an atom with coordinate ri[3] has moved out
206   to neighbor ku.
207 —————*/
208 int bdd(double* ri, int ku) {
209     int kd,kdd;
210     kd = ku/2; /* x(0)/y(1)/z(2) direction */
211     kdd = ku%2; /* Lower(0)/higher(1) direction */
212     if (kdd == 0)
213         return ri[kd] < RCUT;
214     else

```

```

215     return al[kd]-RCUT < ri[kd];
216 }
217 int bmv(double* ri, int ku) {
218     int kd,kdd;
219     kd = ku/2; /* x(0)/y(1)/z(2) direction */
220     kdd = ku%2; /* Lower(0)/higher(1) direction */
221     if (kdd == 0)
222         return ri[kd] < 0.0;
223     else
224         return al[kd] < ri[kd];
225 }

```

pmd_async.c

A.2 Task II

```

1  /*-----
2  Program pmd.c performs parallel molecular-dynamics for Lennard-Jones
3  systems using the Message Passing Interface (MPI) standard.
4  -----*/
5  #include "pmd_async_task2.h"
6  #define VMAX 5.0 // Max. velocity value to construct a velocity
    histogram
7  #define NBIN 100 // # of bins in the histogram
8
9  FILE *fpv;
10
11 /*-----*/
12 int main(int argc, char **argv) {
13 /*-----*/
14     double cpu1,cpu;
15     int i,a;
16
17     MPI_Init(&argc,&argv); /* Initialize the MPI environment */
18     //MPI_Comm_rank(MPI_COMM_WORLD, &sid); /* My processor ID */
19
20     MPI_Comm_rank(MPI_COMM_WORLD, &gid); //Global rank
21     md_shit = gid%2; // = 1 (MD workers) or 0 (analysis workers)
22     MPI_Comm_split(MPI_COMM_WORLD,md_shit,0,&workers);
23     MPI_Comm_rank(workers,&sid); // Rank in workers
24
25     /* Vector index of this processor */
26     vid[0] = sid/(vproc[1]*vproc[2]);
27     vid[1] = (sid/vproc[2])%vproc[1];
28     vid[2] = sid%vproc[2];
29
30     init_params();
31     if (md_shit) {
32         set_topology();

```

```

33     init_conf();
34     atom_copy();
35     compute_accel();
36 }
37 else
38     if (sid == 0)
39         fpv = fopen("pv.dat", "w");
40
41 //printf("%d\n", md_shit);
42 cpu1 = MPI_Wtime();
43 for (stepCount=1; stepCount<=StepLimit; stepCount++) {
44     if (md_shit){
45         printf("%d\n", md_shit);
46         single_step();
47     }
48     if (stepCount%StepAvg == 0) {
49         if (md_shit) {
50             // Send # of atoms, n, to rank gid-1 in MPI_COMM_WORLD
51             MPI_Send(&n, 1, MPI_INT, gid-1, 1000, MPI_COMM_WORLD);
52             // Compose message to be sent
53             for (i=0; i<n; i++)
54                 for (a=0; a<3; a++)
55                     dbuf[3*i+a] = rv[i][a];
56             // Send velocities of n atoms to rank gid-1 in MPI_COMM_WORLD
57             MPI_Send(dbuf, 3*n, MPI_DOUBLE, gid-1, 2000, MPI_COMM_WORLD);
58             eval_props();
59         }
60         else {
61             // Receive # of atoms, n, from rank gid+1 in MPI_COMM_WORLD
62             MPI_Recv(&n, 1, MPI_INT, gid+1, 1000, MPI_COMM_WORLD, &status);
63             // Receive velocities of n atoms from rank gid+1 in
64             // MPI_COMM_WORLD
65             MPI_Recv(dbuf, 3*n, MPI_DOUBLE, gid+1, 2000,
66                     MPI_COMM_WORLD, &status);
67             for (i=0; i<n; i++)
68                 for (a=0; a<3; a++)
69                     rv[i][a] = dbuf[3*i+a];
70             calc_pv();
71         } // end if
72     } // end if
73 } // end for
74
75 cpu = MPI_Wtime() - cpu1;
76 if (md_shit && sid==0)
77     printf("CPU & COMT = %le %le\n", cpu, comt);
78 if (!md_shit && sid==0)
79     fclose(fpv);

```



```
79
80  MPI_Finalize(); /* Clean up the MPI environment */
81  return 0;
82 }
```

pmd_async_task2.c